R at Scale:

Using Apache Spark & Adatao

Christopher Nguyen, PhD
Co-Founder & CEO

Agenda

1. R + Big Data Science: Problem Statement
2. Big Compute: Solution
3. In-Memory Big-Compute: Why & When
4. Apache Spark & Adatao: Overview & Demo

Christopher Nguyen, PhD
Adatao Inc.
Co-Founder & CEO

- Former **Engineering Director of Google Apps** (Google Founders' Award)
- Former Professor and Co-Founder of the **Computer Engineering program at HKUST**
- **PhD Stanford, BS U.C. Berkeley** *Summa cum Laude*
- **Extensive experience** building technology companies that **solve enterprise challenges**

Conventional approach:
Work on sub-sampled data

"We spend 80% of our time shuffling data around.

Parallel computing approach:
Work directly on HDFS

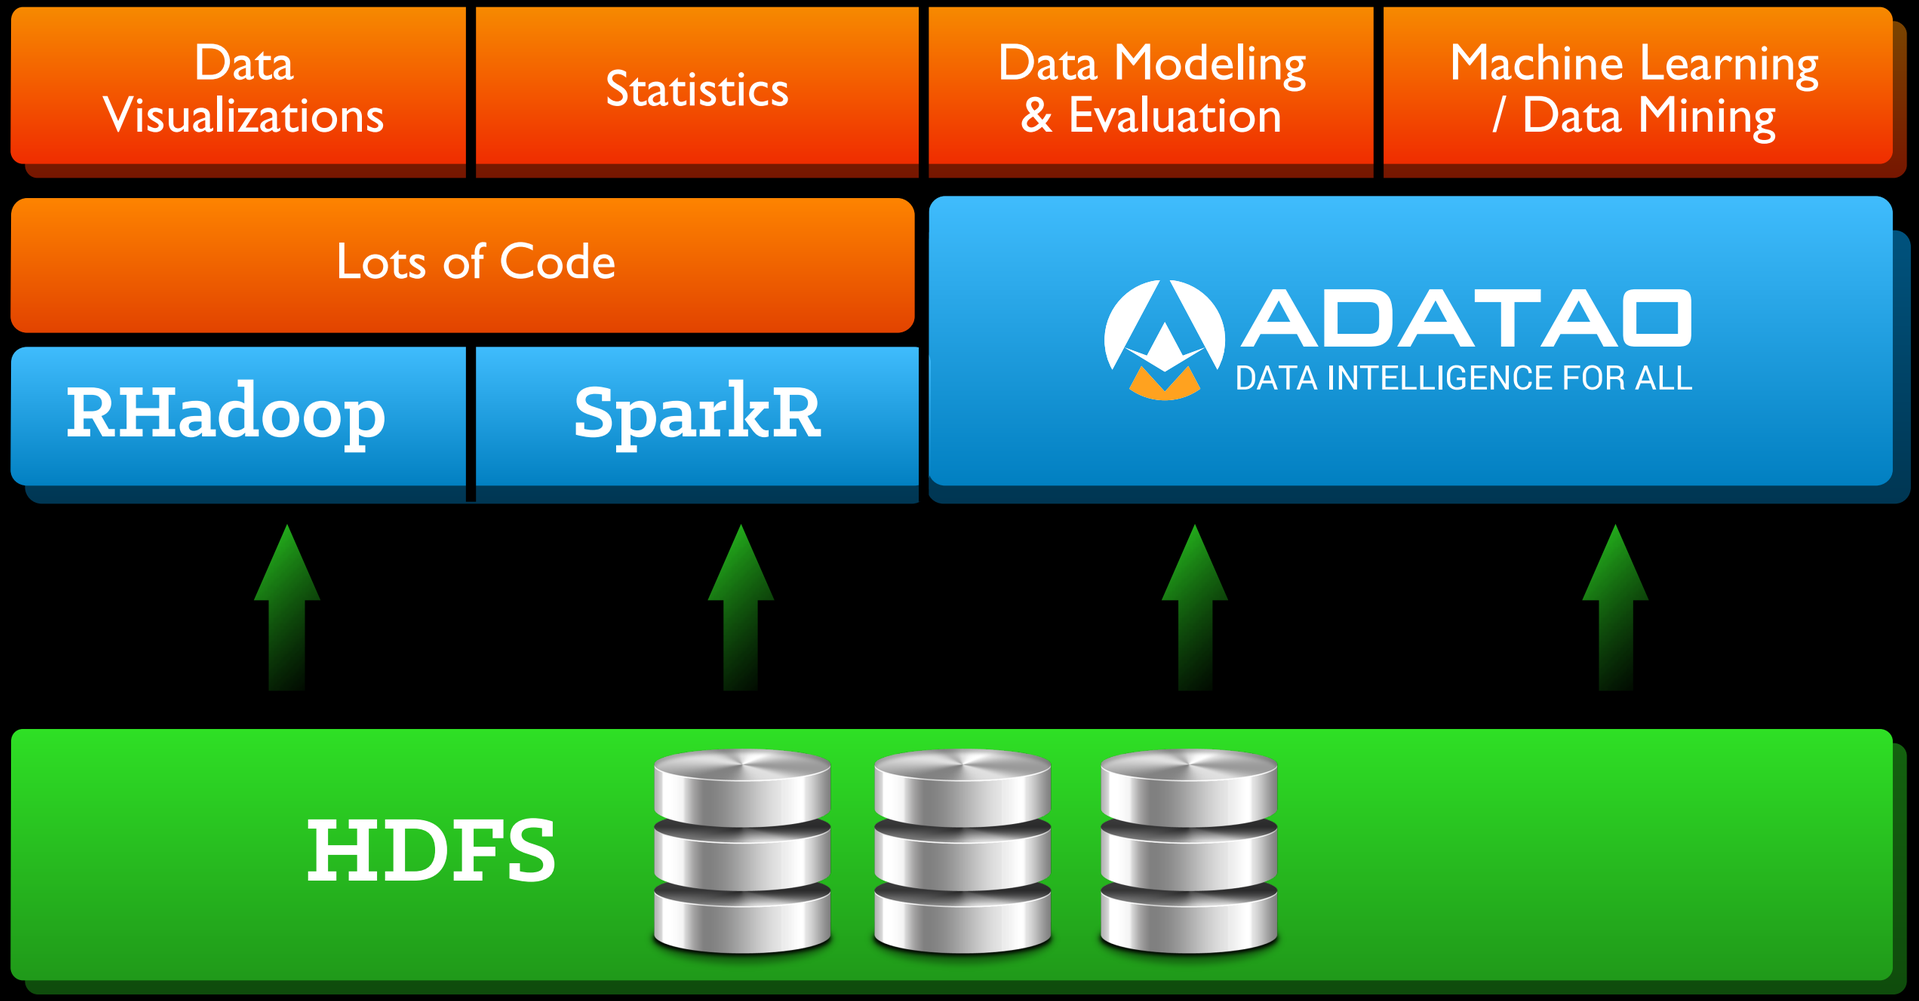

ADATAO

DATA INTELLIGENCE FOR ALL

Big Data & Big Compute

Past & Present

How Have We Defined "Big Data"?

Old Definition

Big Data has Problems

Huge Volume

High Velocity

Great Variety

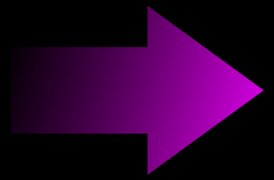

New Definition

BIG DATA + BIG COMPUTE =Opportunitie$

(Machine) Learn from Data

“Big Compute” Defined

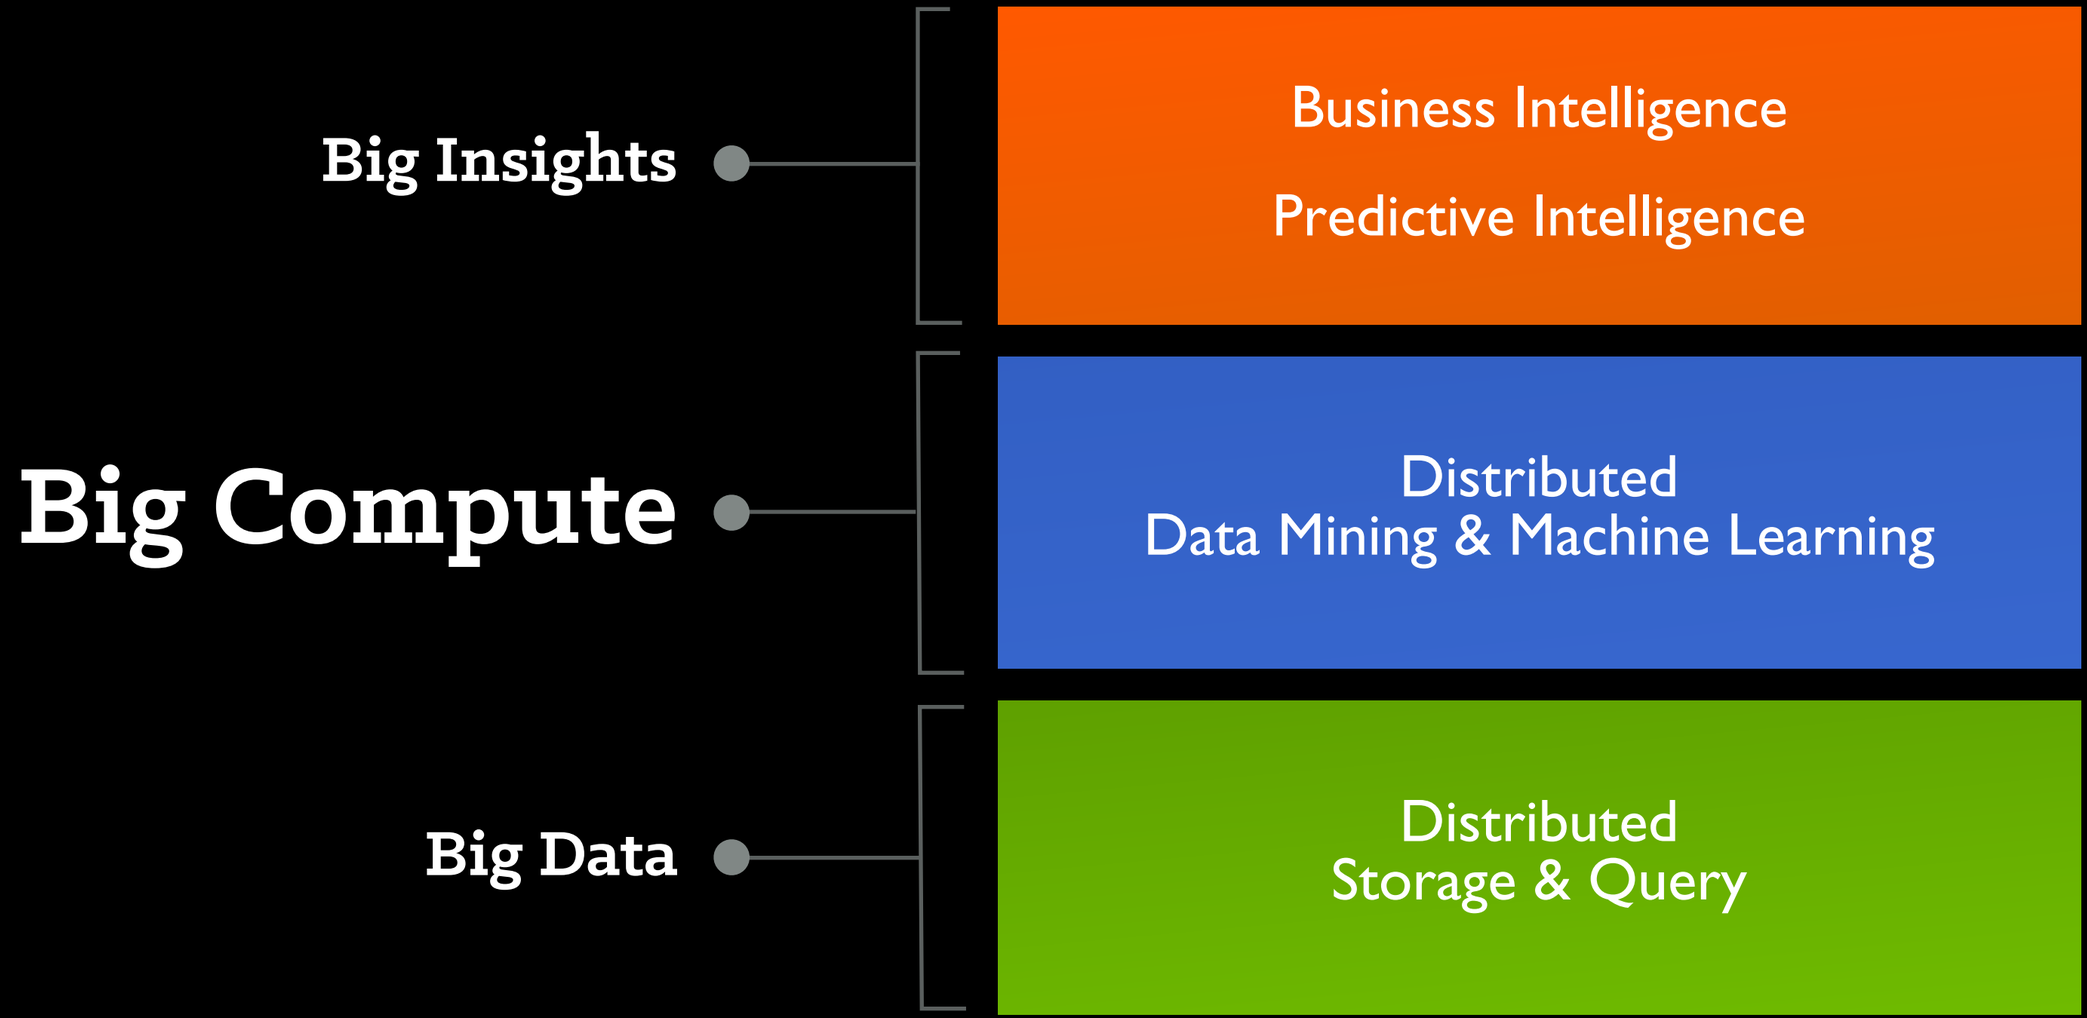

What's Been Missing
In the Big-Data Stack?

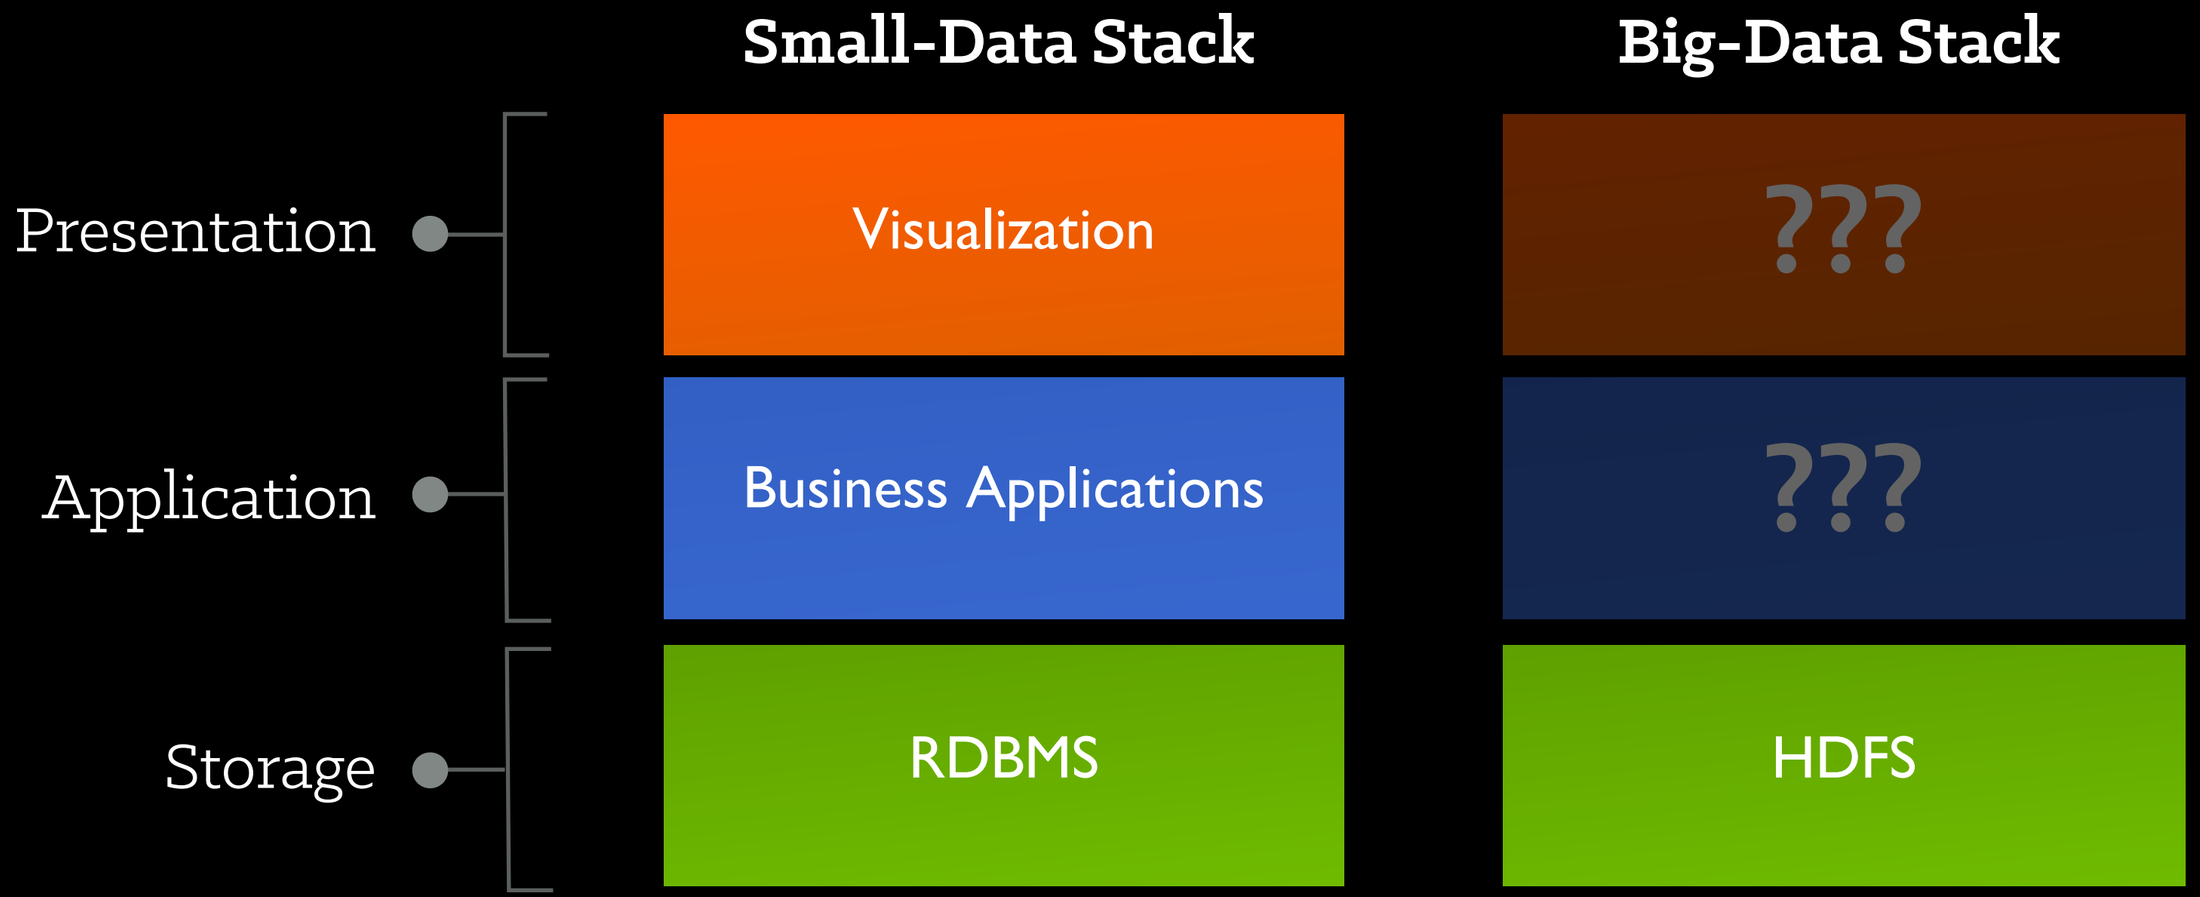

Alphabet Soup

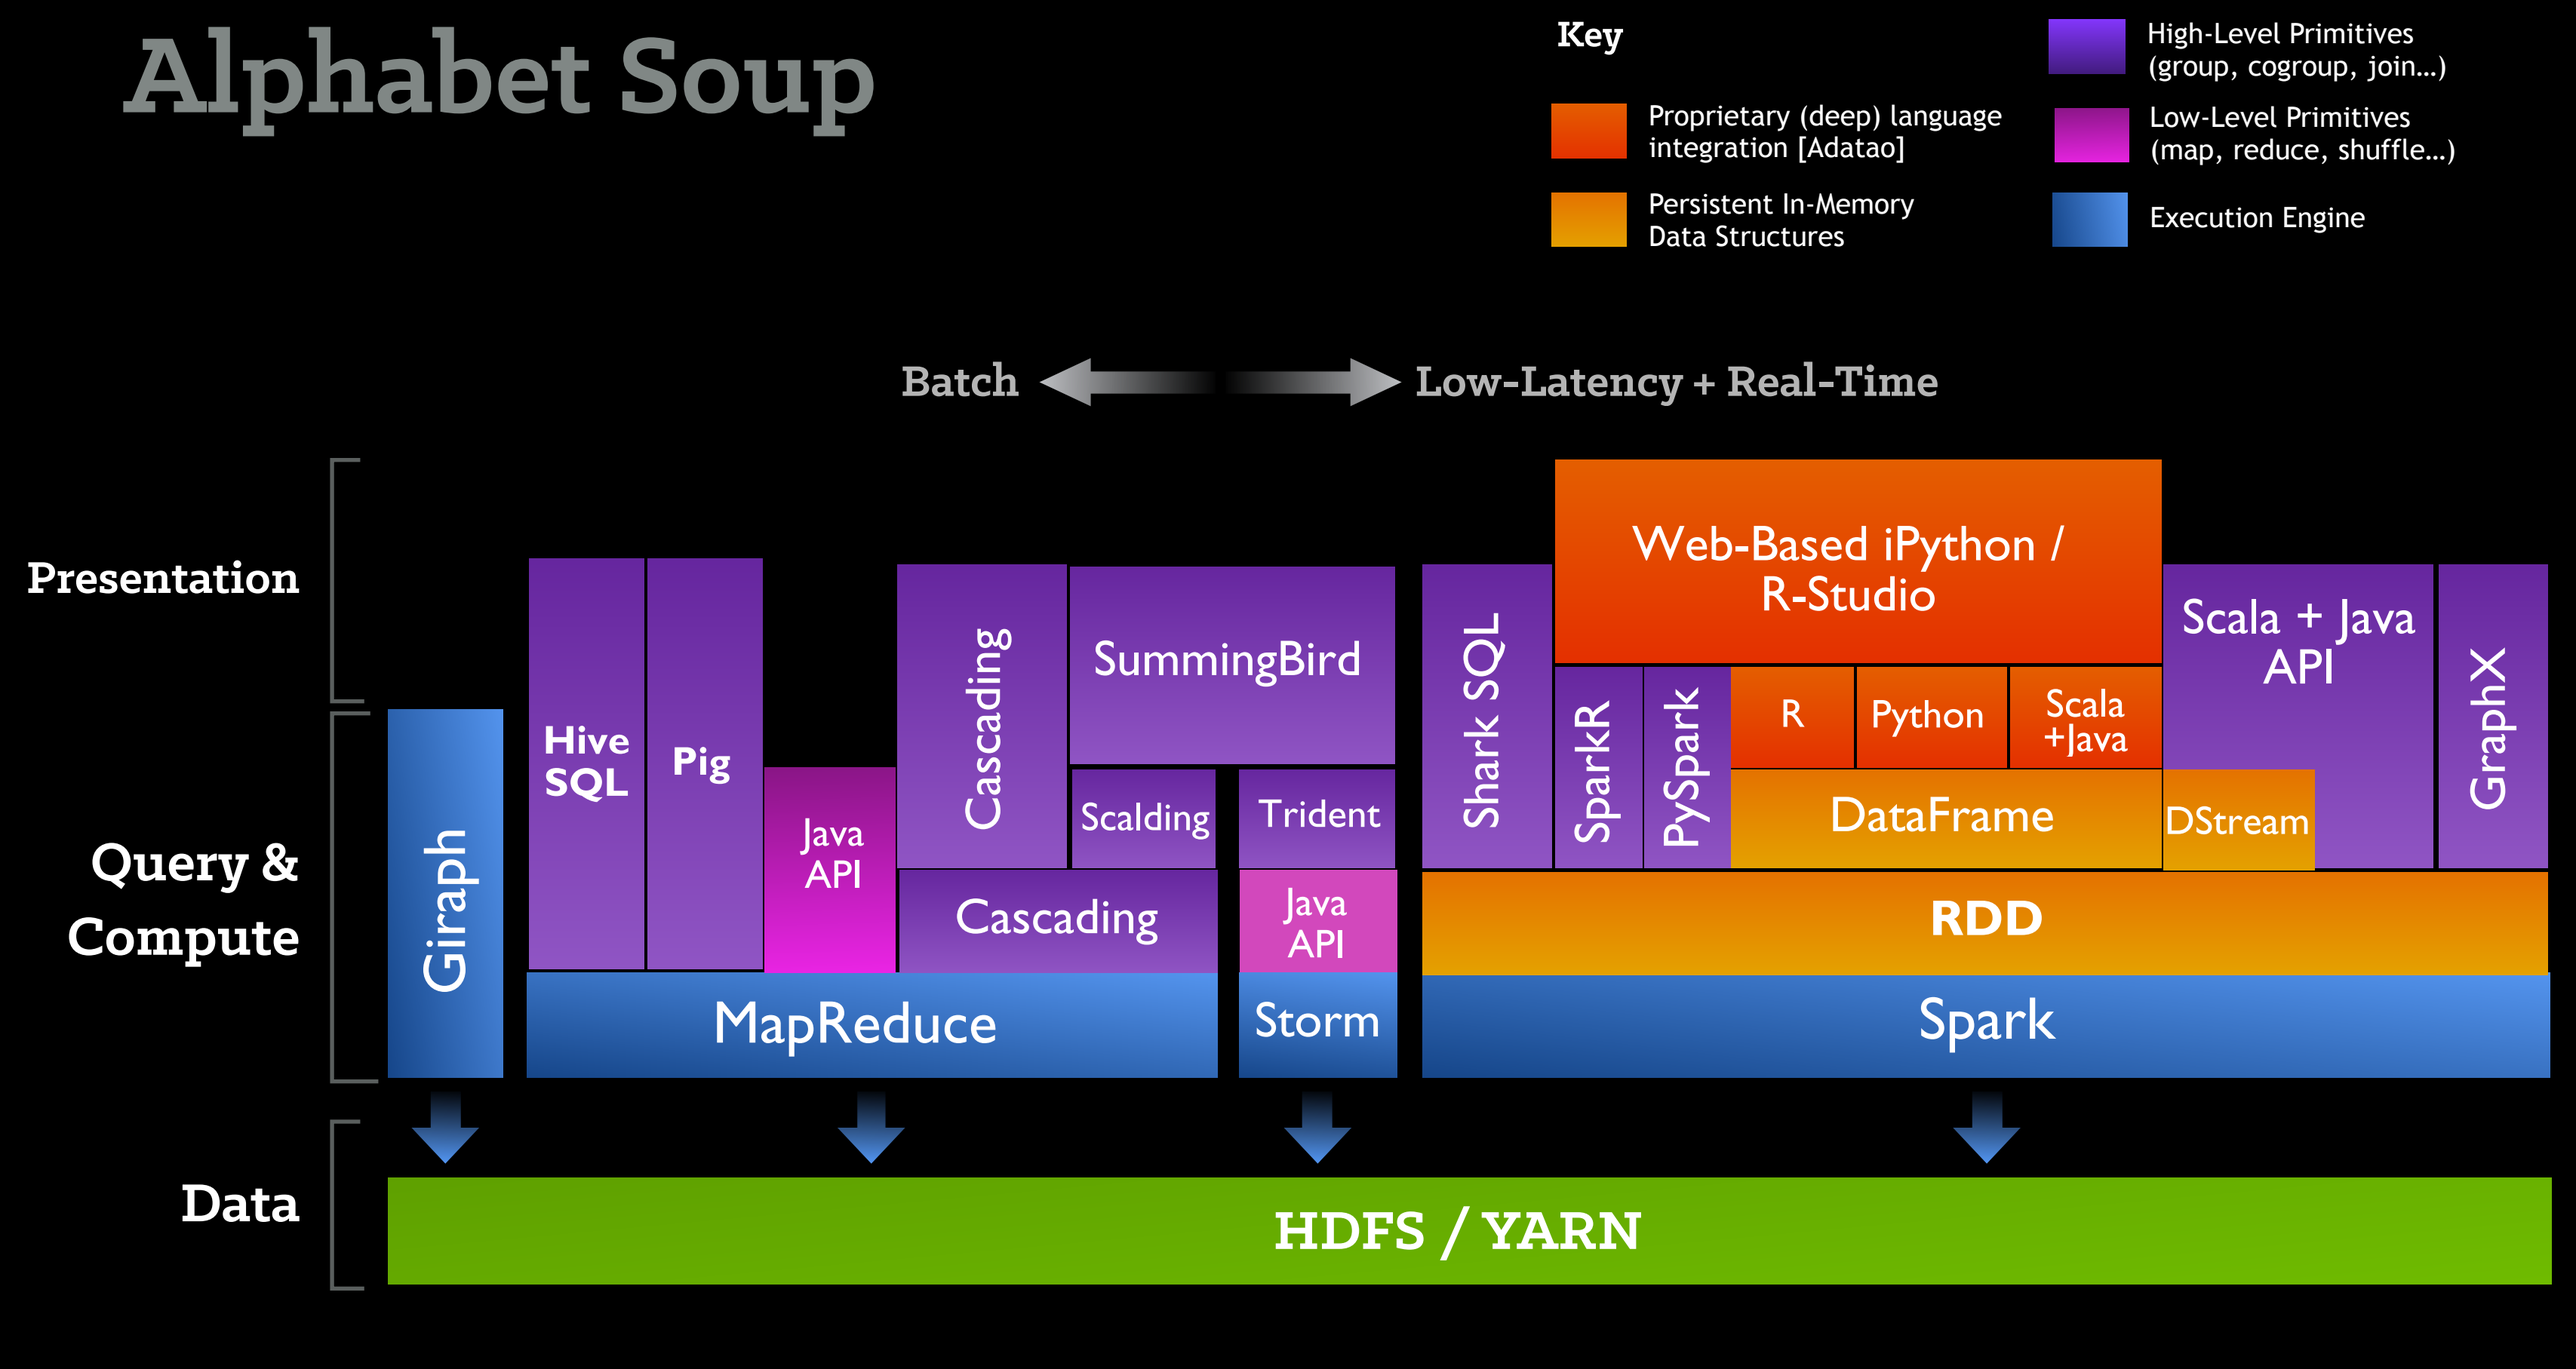

Key components in
the Enterprise Data Flow

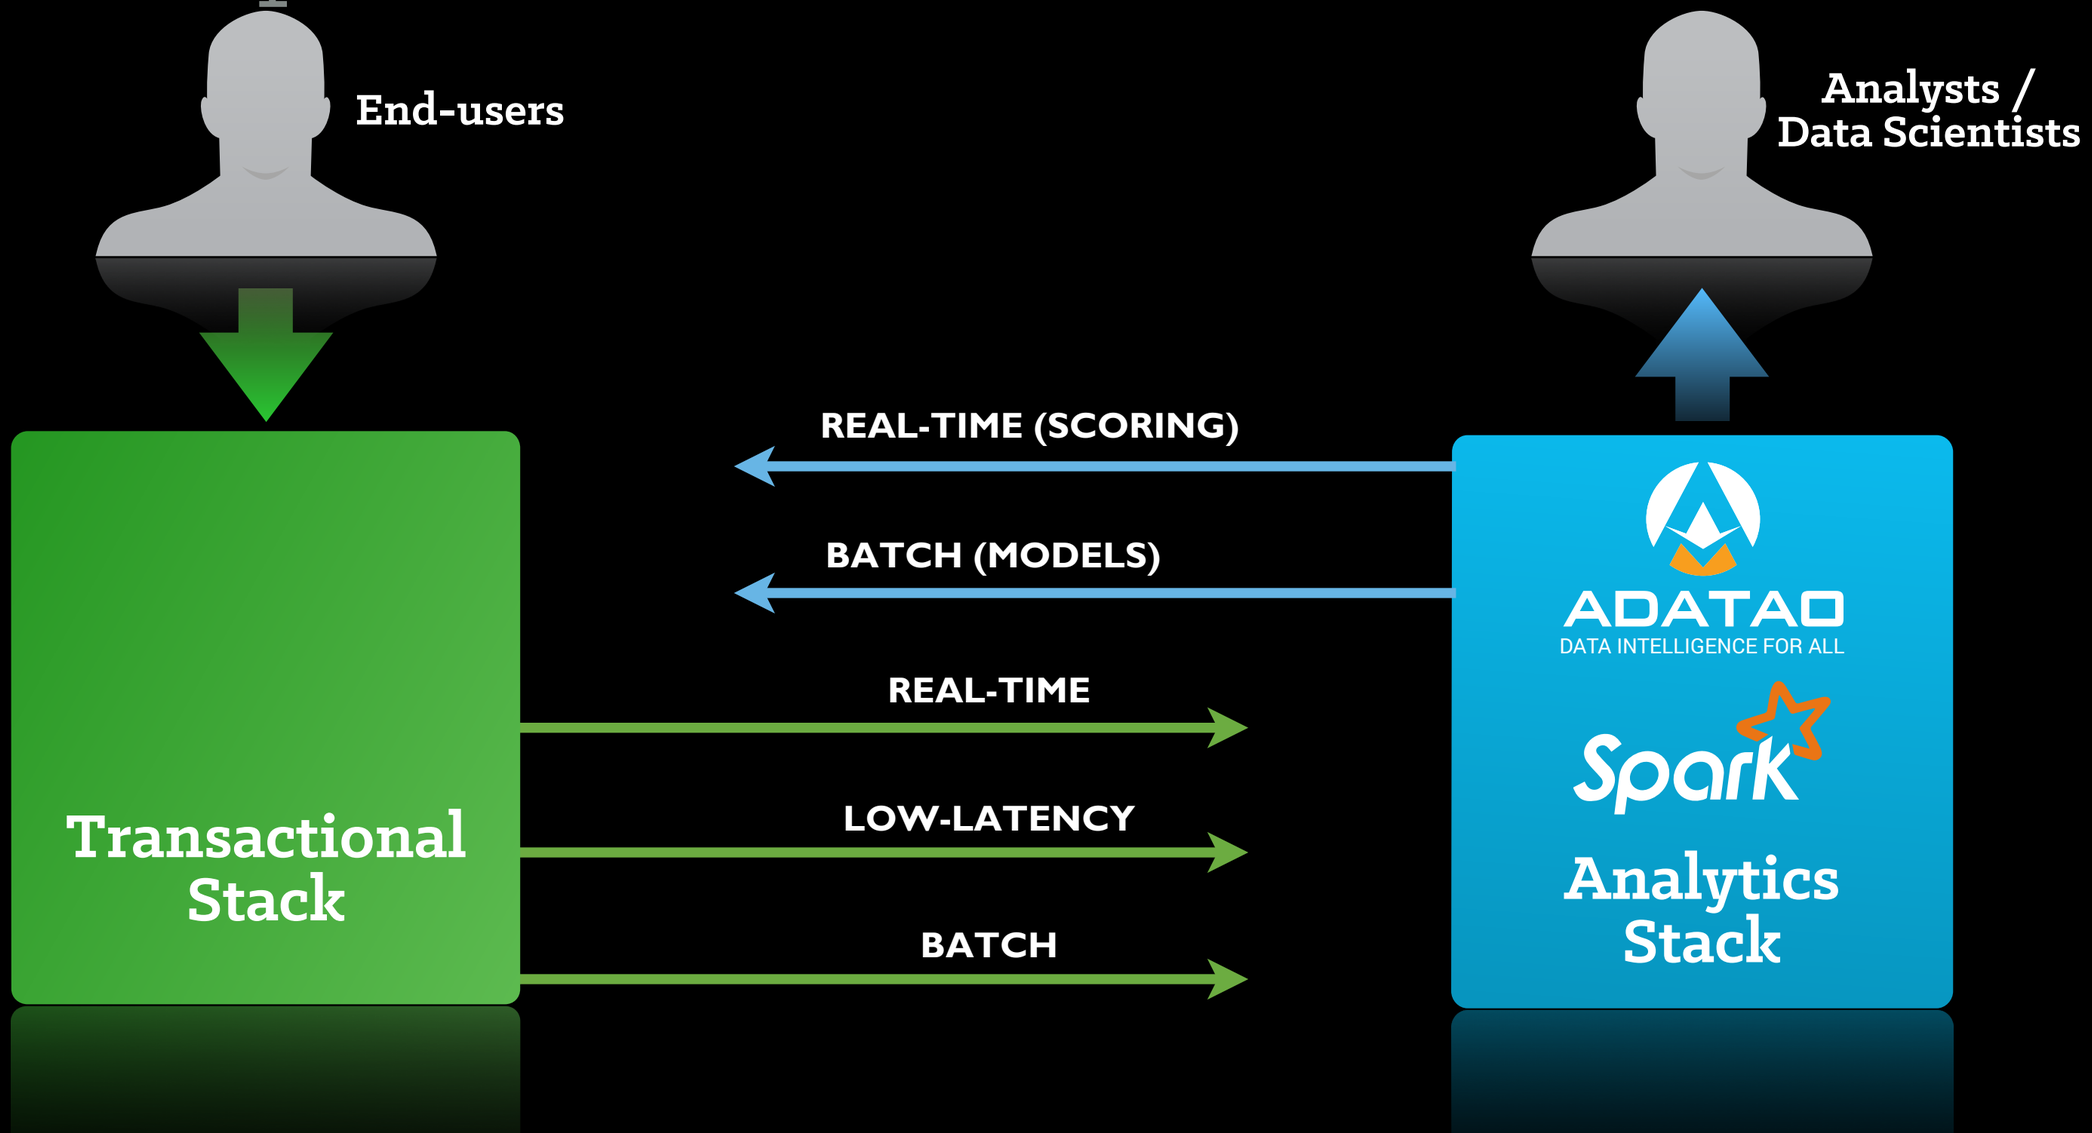

Hadoop
Analytic Landscape

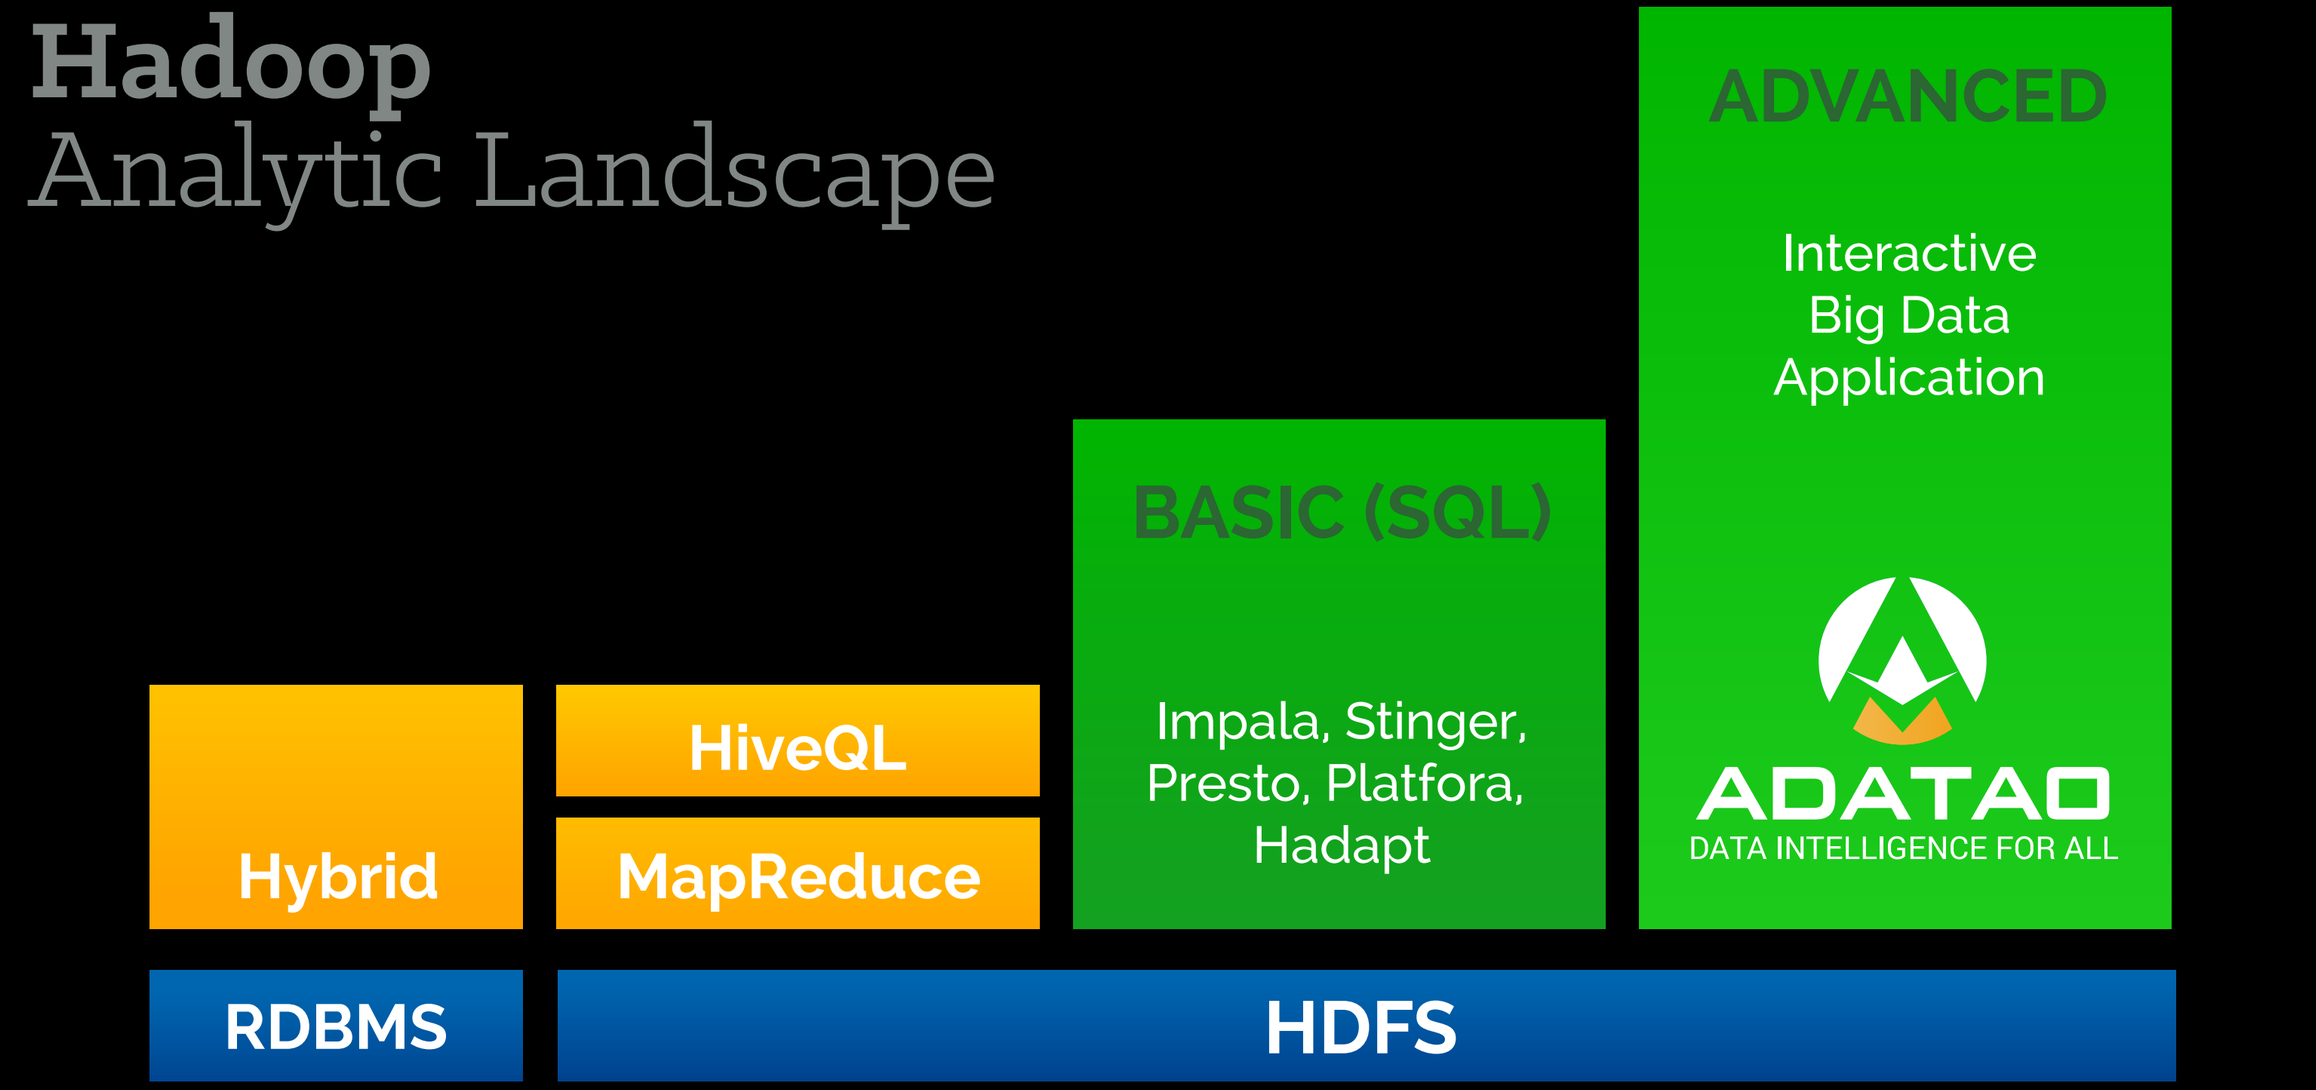

ADATAO

DATA INTELLIGENCE FOR ALL

In-Memory Big Compute

Why? When?

Big-Compute
Value vs Cost Cross-Over Points

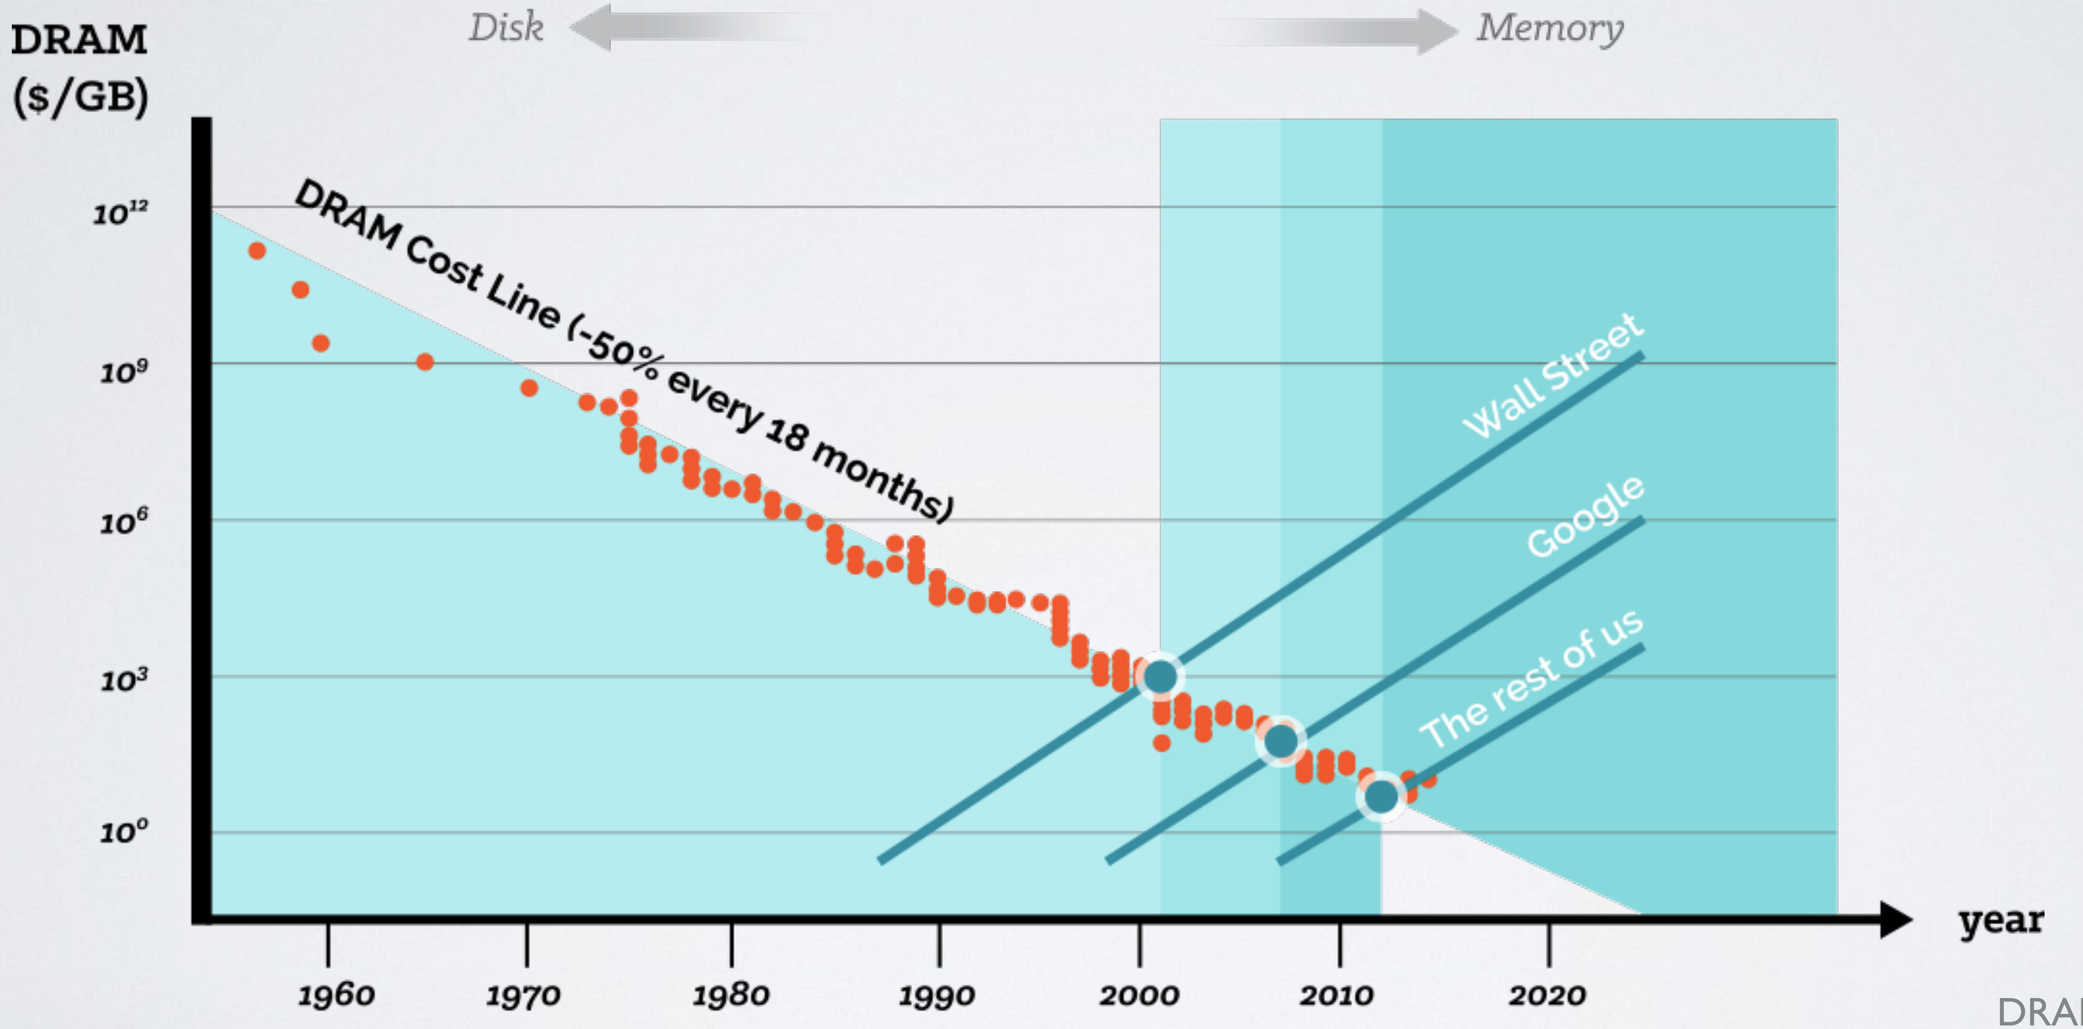

DRAM Source: jcmit.com

The Future Increasingly Favors RAM

Latency Trends

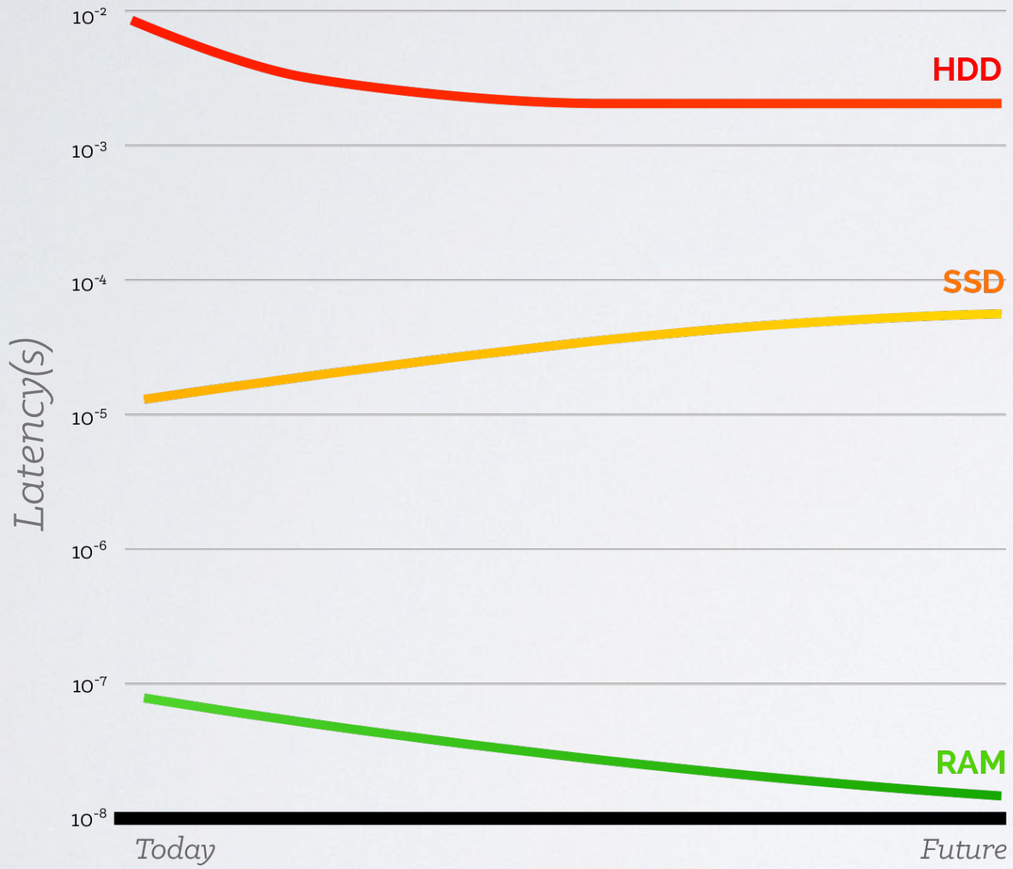

Bandwidth Trends

Unified Big Compute

ADATAO

DATA INTELLIGENCE FOR ALL

Comparison:
Hadoop MapReduce vs. Spark Architecture

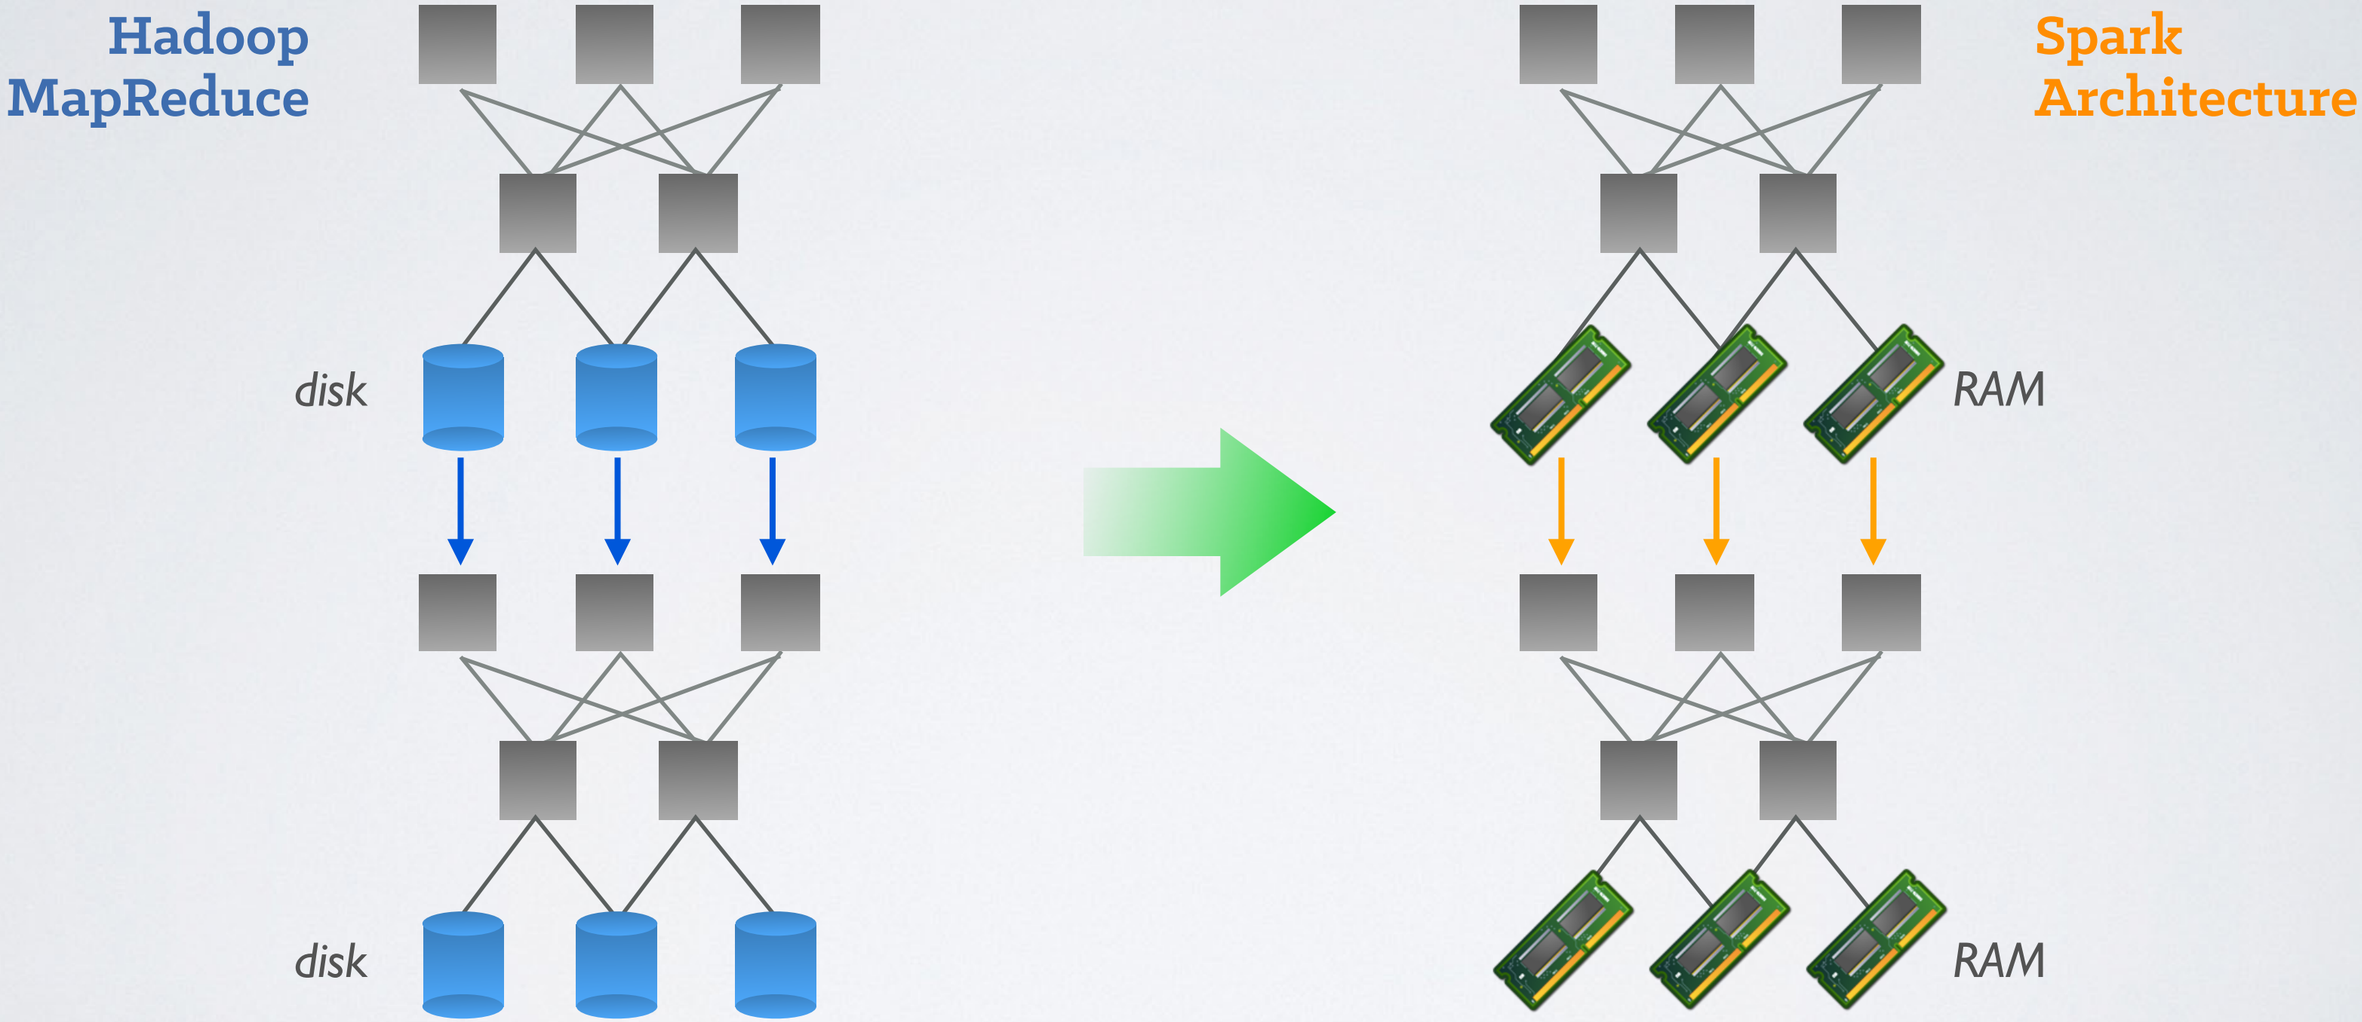

Apache Spark:
Big-Compute Engine

A Compute Engine for Hadoop Data that is:

Fast
Up to **100x** Faster than MapReduce

Sophisticated
Can run today's **most advanced algorithms**

Fully Open Source
One of most active projects in Big Data

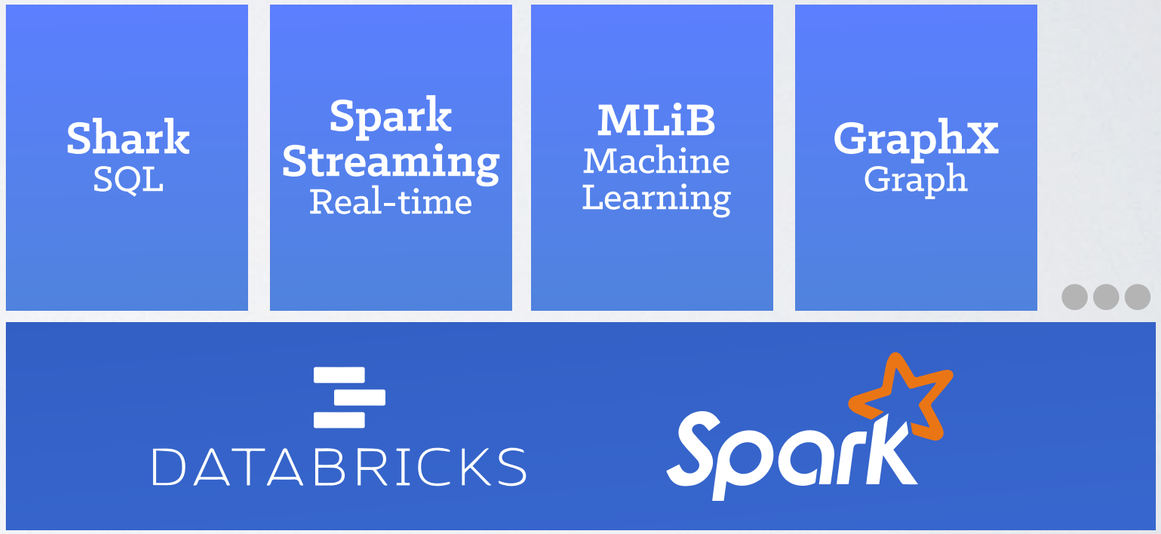

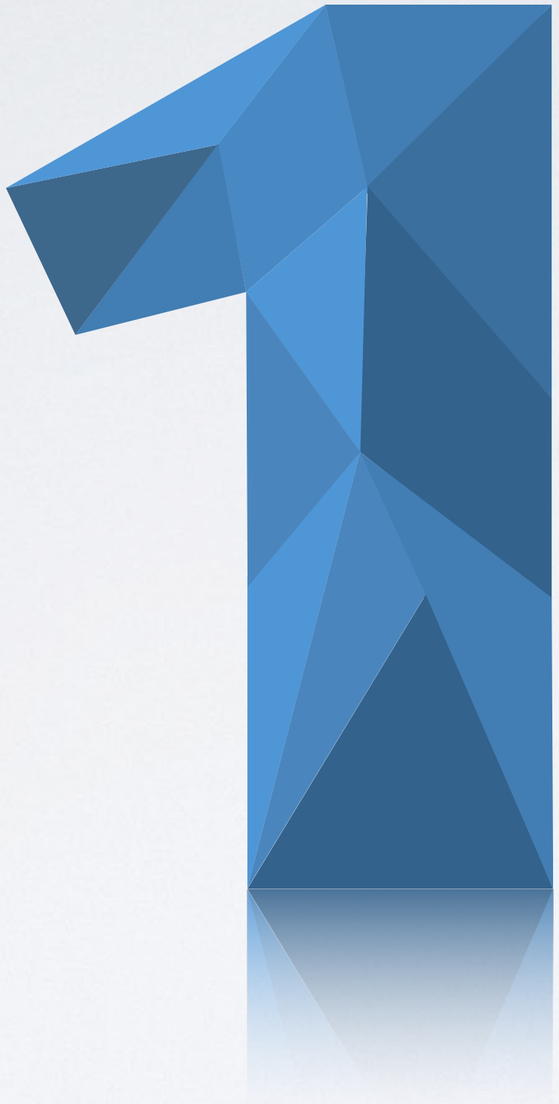

ADATAO

DATA INTELLIGENCE FOR ALL

Unified Workbench *for* Collaborative Data Intelligence

Adatao Architecture

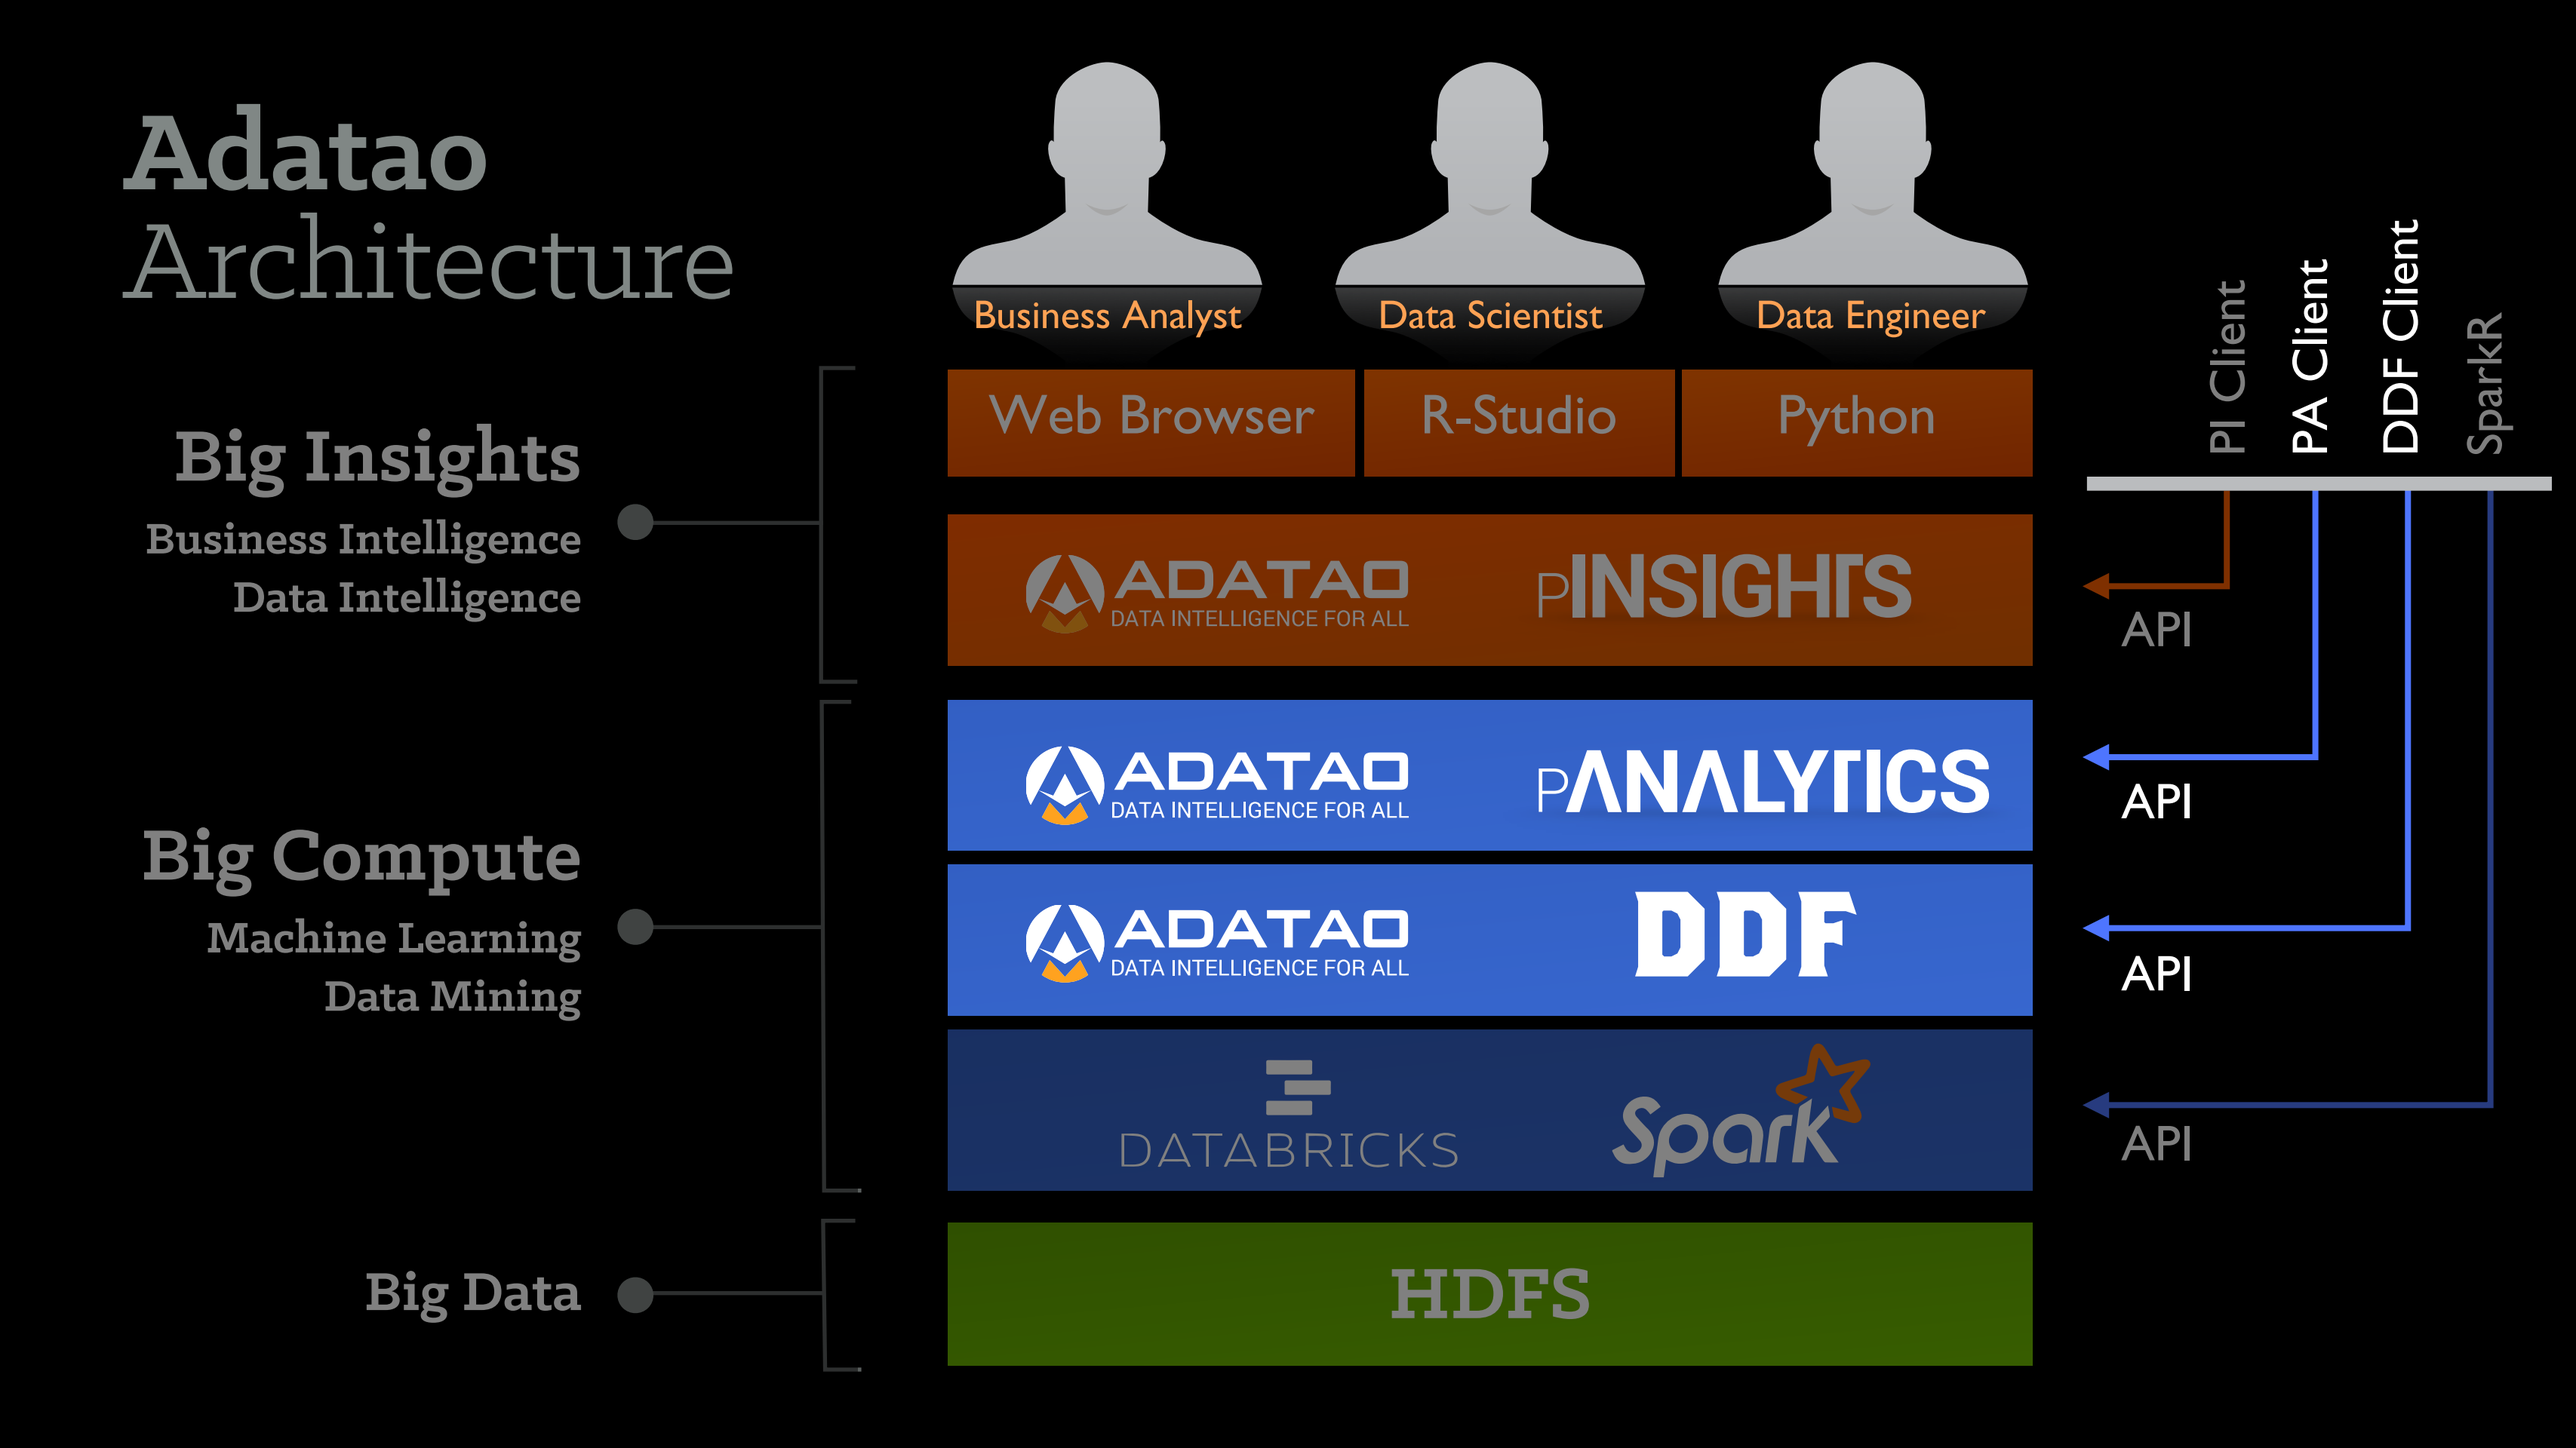

RHadoop

```
library(rmr2)
library(rhdfs)
hdfs.init()
from.dfs(mapreduce(
  input = '/tmp/airline.csv',
  input.format = make.input.format("csv", sep = ","),
  map = function(., data) {
    # filter out non-numeric values (header and NA)
    filter = !is.na(data[,15])
    data = data[filter,]
    # emit composite key (airline|year|month) and delay
    keyval(
    data[,c(9,1,2)],
    data[,15, drop = FALSE])
  },
  reduce = function(k,delays) {
    keyval(k, mean(delays[,1]))
  }
))
```

SparkR

```
library(SparkR)
sc <- sparkR.init()
airlineRDD <- textFile(sc, "/tmp/airline.csv")

map.func <- function(line) {
  data <- unlist(strsplit(line, ","))
  if (data[15] != "NA") { list(data[c(9,1,2)],c(as.integer(data[15]), 1L)) }
}

avg.arrdelay <- lapply(
  reduceByKey(lapply(airlineRDD, map.func),"+",2L),
  function(row) { list(row[[1]], row[[2]][1]/row[[2]][2]) }
)
```

ADATAO
DATA INTELLIGENCE FOR ALL

```
df <- adatao.sql2ddf('select * from airline')
avg.arrdelay <- adatao.aggregate (arrdelay ~ uniquecarrier + year + month, df, FUN=mean)
```

Feature Comparison

	RHadoop	SparkR	Adatao
Support Hive Tables	✗	✗	✓
Support HDFS	✓	✓	✓
Ability to Write MapReduce in R	✓	✓	✓
Native R Idioms	✗	✗	✓
DataFrame Abstraction	✗	✗	✓
Data Extraction	✗	✗	✓
Data Transformation	Raw	Raw	Idiomatic
Data Exploration	✗	✗	✓
Speed	✗	✓	✓✓

Demo Deployment
Diagram

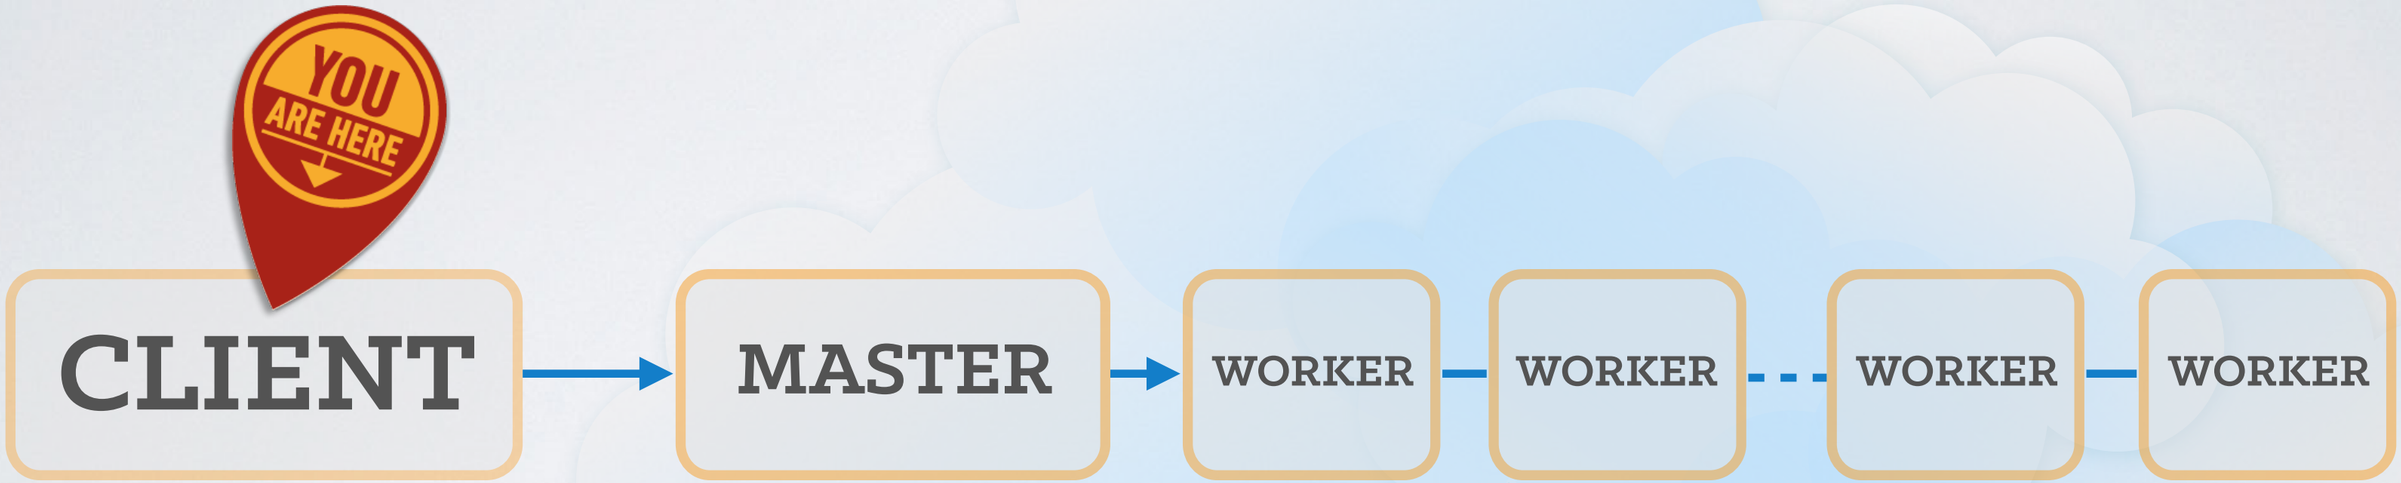

Adatao Benefits

- ✓ **Stop Moving Data Around**
  Data Science Directly on Hadoop Datasets
- ✓ **Native R Data.frame Experience**
  Table-like Abstraction on Top of Big Data
- ✓ **Focus on Analysis, not MapReduce**
  High-Level Programmable API (DDF)
- ✓ **Zero-Effort Model Deployment**
  Transactional & Analytic Support in One Stack
- ✓ **Model Terabytes in Seconds**
  Powerful, Fast, Interactive Data Science
- ✓ **Easily Visualize & Collaborate**
  Beautiful Charting, Dashboarding & RT Collaboration

ADATAO
DATA INTELLIGENCE FOR ALL

Adatao Demo

ADATAO

DATA INTELLIGENCE FOR ALL

To learn more about
Adatao & DDF
contact us, or come to our
Spark Summit talk

www.adatao.com