

Why Are There No Giants?

The Data Analytics of Scalability

Neil J. Gunther @DrQz

Performance Dynamics

Bay Area R Users Group (BARUG) Meetup
September 18, 2018



My vitals



WIKIPEDIA
The Free Encyclopedia



Article [Talk](#)

Neil J. Gunther

From Wikipedia, the free encyclopedia

- Consultant, researcher, author, [teacher](#)
- Ph.D. (U.K.) in theoretical physics
- NASA/JPL [Voyager \(still going!\)](#), [Galileo](#) deep space missions
- Xerox PARC
- Pyramid Technology Corp
- [Performance Dynamics Company](#)
- Using R for past 10 years — practically every day!
- Developed [PDQ \(pretty damn quick\)](#) queueing analyzer — now in R
- Developed the [USL \(universal scalability law\)](#) — R pkgs used in this talk

Wrote some things

[HOME](#)[BLOG](#)[TRAINING](#)[BOOKS](#)[PAPERS](#)[PDQ](#)[GAPHORISMS](#)[SCALABILITY](#)[FACEBOOK](#)[CONTACT](#)

Performance Ponderings


Contents

- [1 Exposing the Cost of Performance Hidden in the Cloud \(2018\)](#)
- [2 WTF is Modeling, Anyway? \(2017\)](#)
- [3 Morphing M/M/m: A New View of an Old Queue \(2017\)](#)
- [4 How to Emulate Web Traffic Using Standard Load Testing Tools \(2016\)](#)
- [5 Hadoop Superlinear Scalability \(2015\)](#)
- [6 A Note on Disk Drag Dynamics \(2012\)](#)
- [7 A Methodology for Optimizing Multithreaded System Scalability on Multicores \(2011\)](#)
- [8 A Note on Parallel Algorithmic Speedup Bounds \(2011\)](#)
- [9 Mind Your Knees and Queues. Responding to Hyperbole with Hyperbole \(2009\)](#)
- [10 A General Theory of Computational Scalability Based on Rational Functions \(2008\)](#)

Outline




- 1 Giant Tales
- 2 Software Scaling
- 3 Universal Scalability Law
- 4 Statistical Regression
- 5 Applying the USL
- 6 Going Further

Mythical giants (and beanstalks)



- J&B giant was reputedly about 30' tall in some accounts (no data)
- Let's not even get into beanstalks in the clouds!

Galileo was onto this




- 1 Discorsi e Dimostrazioni Matematiche Intorno a Due Nuove Scienze¹, Galileo Galilei (1638)
- 2 On Being the Right Size, J. B. S. Haldane (1928)

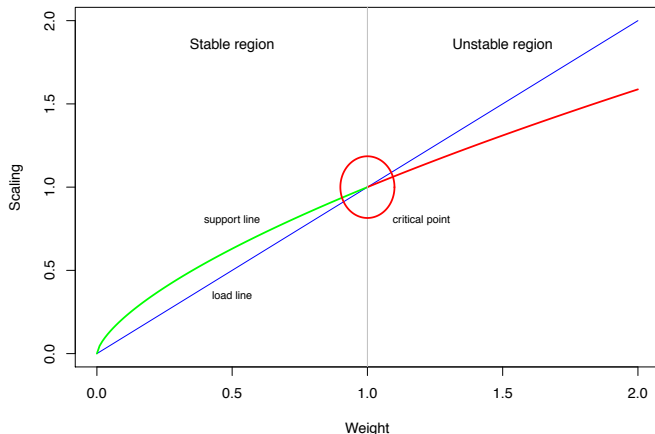
¹ The two new sciences are (i) strength of materials and (ii) projectile motion.

Allometric scaling

- Robert P. Wadlow
- Tallest human giant
- b. Alton, Illinois in 1918
- Reached 8' 11.1" (2.72 meters)
- Guinness world record
- Died in 1940 at 22 yo
- Father was 5' 11.5" (1.82 meters)



Why there are no giants




- 1 Weight (mass) grows like volume: $\sim L^3$
- 2 Support only grows like cross-sectional area: $\sim L^2$


Outline

- 1 Giant Tales
- 2 Software Scaling
- 3 Universal Scalability Law
- 4 Statistical Regression
- 5 Applying the USL
- 6 Going Further


Equal bang for your buck: Concurrency




Diminishing returns: Contention cost



Resource saturation: More contention cost



Negative returns: Coherency of distributed data



Outline

- 1 Giant Tales
- 2 Software Scaling
- 3 Universal Scalability Law**
- 4 Statistical Regression
- 5 Applying the USL
- 6 Going Further

Model: Universal scalability law (USL)

N : processes provide system stimulus or **load**

²NJG. "A Simple Capacity Model of Massively Parallel Transaction Systems," *CMG Conf.* (1993)

Model: Universal scalability law (USL)

N : processes provide system stimulus or **load**

C_N : response function or **relative capacity**

²NJG. "A Simple Capacity Model of Massively Parallel Transaction Systems," *CMG Conf.* (1993)

Model: Universal scalability law (USL)

N : processes provide system stimulus or **load**

C_N : response function or **relative capacity**

Question: What kind of function?

²NJG. "A Simple Capacity Model of Massively Parallel Transaction Systems," *CMG Conf.* (1993)

Model: Universal scalability law (USL)

N : processes provide system stimulus or **load**

C_N : response function or **relative capacity**

Question: What kind of function?

Answer: Nonlinear rational function²

²NJG. "A Simple Capacity Model of Massively Parallel Transaction Systems," *CMG Conf.* (1993)

Model: Universal scalability law (USL)

N : processes provide system stimulus or **load**

C_N : response function or **relative capacity**

Question: What kind of function?

Answer: Nonlinear **rational function**²

$$C_N(\alpha, \beta, \gamma) = \frac{\gamma N}{1 + \alpha(N-1) + \beta N(N-1)}$$

²NJG. "A Simple Capacity Model of Massively Parallel Transaction Systems," *CMG Conf.* (1993)

Model: Universal scalability law (USL)

N : processes provide system stimulus or **load**

C_N : response function or **relative capacity**

Question: What kind of function?

Answer: Nonlinear **rational function**²

$$C_N(\alpha, \beta, \gamma) = \frac{\gamma N}{1 + \alpha(N-1) + \beta N(N-1)}$$

The three Cs:

²NJG. "A Simple Capacity Model of Massively Parallel Transaction Systems," *CMG Conf.* (1993)

Model: Universal scalability law (USL)

N : processes provide system stimulus or **load**

C_N : response function or **relative capacity**

Question: What kind of function?

Answer: Nonlinear **rational function**²

$$C_N(\alpha, \beta, \gamma) = \frac{\gamma N}{1 + \alpha(N-1) + \beta N(N-1)}$$

The three Cs:

① **C**oncurrency ($1 < \gamma < \infty$)

²NJG. "A Simple Capacity Model of Massively Parallel Transaction Systems," *CMG Conf.* (1993)

Model: Universal scalability law (USL)

N : processes provide system stimulus or **load**

C_N : response function or **relative capacity**

Question: What kind of function?

Answer: Nonlinear **rational function**²

$$C_N(\alpha, \beta, \gamma) = \frac{\gamma N}{1 + \alpha(N-1) + \beta N(N-1)}$$

The three Cs:

① **C**oncurrency ($1 < \gamma < \infty$)

② **C**ontention ($0 < \alpha < 1$)

²NJG. "A Simple Capacity Model of Massively Parallel Transaction Systems," *CMG Conf.* (1993)

Model: Universal scalability law (USL)

N : processes provide system stimulus or **load**

C_N : response function or **relative capacity**

Question: What kind of function?

Answer: Nonlinear **rational function**²

$$C_N(\alpha, \beta, \gamma) = \frac{\gamma N}{1 + \alpha(N-1) + \beta N(N-1)}$$

The three Cs:

- ① **C**oncurrency ($1 < \gamma < \infty$)
- ② **C**ontention ($0 < \alpha < 1$)
- ③ **C**oherency ($0 < \beta < \infty$)

²NJG. "A Simple Capacity Model of Massively Parallel Transaction Systems," *CMG Conf.* (1993)

Measurement meets model

$$\underbrace{\frac{X(N)}{X(1)}}_{\text{Thruput data}}$$

Measurement meets model


$$\underbrace{\frac{X(N)}{X(1)}}_{\text{Thruput data}} \longrightarrow \underbrace{C_N}_{\text{Scalability}}$$

Measurement meets model

$$\underbrace{\frac{X(N)}{X(1)}}_{\text{Thruput data}} \longrightarrow \underbrace{C_N}_{\text{Scalability}} \longleftarrow \underbrace{\frac{\gamma N}{1 + \alpha (N-1) + \beta N(N-1)}}_{\text{USL model}}$$

Measurement meets model

$$\underbrace{\frac{X(N)}{X(1)}}_{\text{Thruput data}} \rightarrow \underbrace{C_N}_{\text{Scalability}} \leftarrow \underbrace{\frac{\gamma N}{1 + \alpha (N-1) + \beta N(N-1)}}_{\text{USL model}}$$



Outline

- 1 Giant Tales
- 2 Software Scaling
- 3 Universal Scalability Law
- 4 Statistical Regression**
- 5 Applying the USL
- 6 Going Further

Determining α, β, γ coefficients

$$C_N(\alpha, \beta, \gamma) = \frac{\gamma N}{1 + \alpha (N - 1) + \beta N(N - 1)}$$

- Brute force measurements (good luck!)

Determining α, β, γ coefficients

$$C_N(\alpha, \beta, \gamma) = \frac{\gamma N}{1 + \alpha (N - 1) + \beta N(N - 1)}$$


- Brute force measurements (good luck!)
- **Clever way:** Apply statistical regression

Determining α, β, γ coefficients

$$C_N(\alpha, \beta, \gamma) = \frac{\gamma^N}{1 + \alpha(N-1) + \beta N(N-1)}$$


- Brute force measurements (good luck!)
- **Clever way:** Apply [statistical regression](#)
- Magic functions in **R**:
 - `nls()` **nonlinear regression** → α, β, γ in one swell foop
 - `predict()` smooth interpolation/extrapolation from data
 - `plot()` cleanest for iterative assessment
 - `ggplot()` with all its bells and whistles

Least squares history




NASA [Dawn spacecraft](#) is currently orbiting dwarf planet Ceres

Least squares history




- Ceres was first observed more than 200 years ago, viz., 1801
- Mathematician C.F. Gauss most accurately estimated the (then unknown) orbit of Ceres
- He developed **least squares** statistical regression for this purpose
- Data errors represented by Gaussian (“normal”) distribution

Those squares are real



Calculated by “linear model” `lm()` base function in R

Coefficient of determination




Multiple R-squared: 0.7494 from $R^2 = 1 - SSE/SST$

Outline

- 1 Giant Tales
- 2 Software Scaling
- 3 Universal Scalability Law
- 4 Statistical Regression
- 5 Applying the USL**
- 6 Going Further

Memcache scalability data



Example (Nonlinear regression)

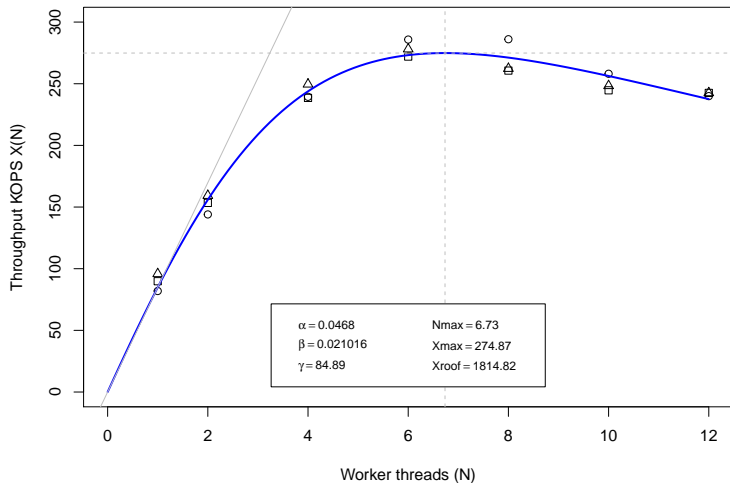
```
# plot the X(N) in KOPS against active worker threads N
plot(df.mc$N,df.mc$mc128,
     xlim=c(0,12), ylim=c(0,300),
     xlab="Worker threads (N)",ylab="Throughput KOPS X(N)"
)
points(df.mc$N,df.mc$mc141,pch=0)
points(df.mc$N,df.mc$mc145,pch=2)

usl.fit <- nls(X.mean ~ gamma * N /
              (1 + alpha * (N - 1) + beta * N * (N - 1)),
              data=df.mc, start=list(gamma=100, alpha=1e-2, beta=1e-2),
              algorithm="port",
              lower=c(50,0,0), upper=c(200,5e-1,5e-1)
              )

# overlay continuous USL curve in blue
curve(coef(usl.fit)['gamma'] * x /
      (1 + coef(usl.fit)['alpha'] * (x - 1) +
       coef(usl.fit)['beta'] * x * (x - 1)),
      from=0, to=12, add=TRUE, col="blue", lwd=2
      )



# compute USL scaling metrics
Nmax <- sqrt((1 - coef(usl.fit)['alpha']) / coef(usl.fit)['beta'])
Xmax <- coef(usl.fit)['gamma'] * Nmax /
      (1 + coef(usl.fit)['alpha'] * (Nmax - 1) +
       coef(usl.fit)['beta'] * Nmax * (Nmax - 1))
Nopt <- abs(1 / coef(usl.fit)['alpha'])
Xroof <- coef(usl.fit)['gamma'] * Nopt
```

Memcache scalability model



MySQL big data

Analyzed some 2,000 production database logs
I resorted to animation as a visualization tool
About 500,000 data points in aggregate



Comparison of USL parameters
USL found unexpected progressive changes in scalability

R packages for the USL

1 SATK on R-Forge

- Author: Paul Puglia (Guerrilla graduate)
- Applies multiple USL coefficient models for best fit

```
install.packages("SATK", repos="http://R-Forge.R-project.org")
```

```
library(SATK)
data(USLcalc)
uslcalc.zones <- zones(USLcalc)
plot(uslcalc.zones)
```

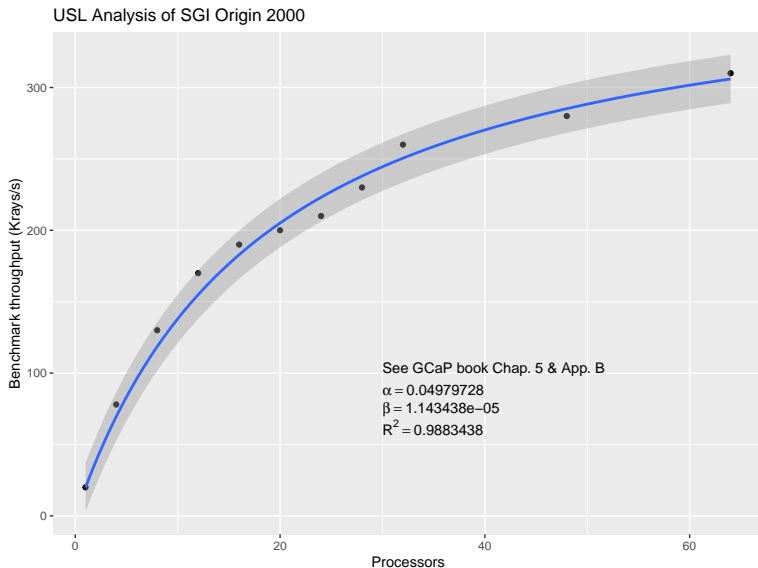
2 usl on CRAN

- Author: Stefan Möding
- Uses both `nls()` from base and `nlxb()` from [nlxr package](#)

```
install.packages("usl")
```

```
library(usl)
data(specsdm91)
usl.model <- usl(throughput ~ load, specsdm91)
summary(usl.model)
peak.scalability(usl.model)
plot(specsdm91, pch=16)
plot(usl.model, col="red", add=TRUE)
```


Applying ggplot enhancements



Outline

- 1 Giant Tales
- 2 Software Scaling
- 3 Universal Scalability Law
- 4 Statistical Regression
- 5 Applying the USL
- 6 Going Further**


Wanna learn more?



- Chap. 4 — Scalability: A Quantitative Approach
- Chap. 5 — Evaluating Scalability Parameters
- Chap. 6 — Software Scalability

- [Guerrilla training classes](#) — textbooks included
- [Guerrilla Data Analytics](#) class — linear regression to machine learning in R

A “darker” universal law?



Questions?

www.perfdynamics.com
Castro Valley, California

Twitter
Facebook
Blog
Training classes
info@perfdynamics.com
+1-510-537-5758